



BSDS-209 - NLP for Data Science

Course Instructor - Mr. Nitish Patil

Gauri Sharan - BSc Data Science (Sem 3)

January 6, 2024

Outline

- 1 Overview
- 2 Applications of NLP
- 3 NLP Pipeline
- 4 Tokenization in Python
- 5 Regex in Python
- 6 Lemmatization in Python
- 7 Stemming in Python
- 8 POS Tagging in Python
- 9 Stopwords Removal in Python
- 10 One-Hot Encoding in Python
- 11 Bag of Words in Python
- 12 TF-IDF in Python
- 13 N-grams in Python
- 14 Named Entity Recognition (NER) in Python

What is NLP?

- Natural Language Processing (NLP) is a subfield of artificial intelligence (AI) that focuses on the interaction between computers and humans using natural language.
- It involves the development of algorithms and models to enable computers to understand, interpret, and generate human language.

Applications of NLP

- Sentiment Analysis
- Named Entity Recognition
- Machine Translation
- Text Summarization
- Speech Recognition
- Chatbots and Virtual Assistants

NLP Pipeline

- 1 Data Acquisition
- 2 Text Preprocessing
 - Text Cleaning - HTML tag removing, emoji handling, Spelling checker, etc.
 - Basic Preprocessing - (word tokenization, sentence tokenization, stop word removal)
 - Advance Preprocessing - POS tagging, Parsing, etc.
- 3 Featured Engineering
 - One-Hot Encoder
 - Bag-of-Words
 - n-grams
 - tf-idf
 - word2vec
- 4 Model Building
- 5 Evaluation and Deployment

Tokenization in Python

```
import nltk
from nltk.tokenize import word_tokenize

text = "Tokenization-is-the-process-of-breaking-down
text-into-words."
tokens = word_tokenize(text)
print(tokens)
```

```
['Tokenization', 'is', 'the', 'process', 'of', 'breaking', 'down', 'text', 'into', 'words', '.']
```

Regex in Python

```
import re
```

```
text = "Regular-expressions-are-powerful-for-text  
pattern-matching."
```

```
pattern = r'\b\w+\b' # Matches words
```

```
matches = re.findall(pattern, text)
```

```
print(matches)
```

```
['Regular', 'expressions', 'are', 'powerful', 'for', 'text', 'pattern', 'matchi  
ng']
```

Lemmatization in Python

```
from nltk.stem import WordNetLemmatizer

lemmatizer = WordNetLemmatizer()
words = ["running", "better", "cats"]
lemmatized_words = [lemmatizer.lemmatize(word) for
word in words]
print(lemmatized_words)
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
['running', 'better', 'cat']
```


Stemming in Python

```
from nltk.stem import PorterStemmer
```

```
stemmer = PorterStemmer()
```

```
words = ["running", "better", "cats"]
```

```
stemmed_words = [stemmer.stem(word) for word in words]
```

```
print(stemmed_words)
```

```
print(stemmed_words)
```

```
['run', 'better', 'cat']
```

POS Tagging in Python

```
import nltk
```

```
text = "POS-tagging-assigns-parts-of-speech-to-words  
in-a-sentence."
```

```
pos_tags = nltk.pos_tag(nltk.word_tokenize(text))  
print(pos_tags)
```

```
[('POS', 'NNP'),  
 ('tagging', 'VBG'),  
 ('assigns', 'JJ'),  
 ('parts', 'NNS'),  
 ('of', 'IN'),  
 ('speech', 'NN'),  
 ('to', 'TO'),  
 ('words', 'NNS'),  
 ('in', 'IN'),  
 ('a', 'DT'),  
 ('sentence', 'NN'),  
 ('.', '.')]
```

Stopwords Removal in Python

```
from nltk.corpus import stopwords

text = "Remove stopwords from this text."
stopwords_set = set(stopwords.words('english'))
filtered_words = [word for word in
nltk.word_tokenize(text) if word.lower() not in
stopwords_set]
print(filtered_words)
```

```
['Remove', 'stopwords', 'text', '.']
```

One-Hot Encoding in Python

```
from sklearn.preprocessing import OneHotEncoder

categories = ['category_A', 'category_B',
             'category_C']
encoder = OneHotEncoder()
encoded_data = encoder.fit_transform([[cat] for cat
in categories]).toarray()
print(encoded_data)
```

```
[[1.  0.  0.]
 [0.  1.  0.]
 [0.  0.  1.]]
```

Bag of Words in Python

```
from sklearn.feature_extraction.text import  
CountVectorizer
```

```
corpus = ['This-is-the-first-document.', 'This-is-the  
second-document.']  
vectorizer = CountVectorizer()  
bag_of_words =  
vectorizer.fit_transform(corpus).toarray()  
print(bag_of_words)
```

```
print(bag_of_words)  
  
[[1 1 1 0 1 1]  
 [1 0 1 1 1 1]]
```

TF-IDF in Python

```
from sklearn.feature_extraction.text import  
TfidfVectorizer
```

```
corpus = ['This-is-the-first-document.', 'This  
is-the-second-document.']  
vectorizer = TfidfVectorizer()  
tfidf_matrix =  
vectorizer.fit_transform(corpus).toarray()  
print(tfidf_matrix)
```

```
[[0.4090901  0.57496187 0.4090901  0.          0.4090901  0.4090901 ]  
 [0.4090901  0.          0.4090901  0.57496187 0.4090901  0.4090901 ]]
```

N-grams in Python

```
from nltk import ngrams
```

```
sentence = "Generate n-grams from this sentence."  
n = 2 # Change to desired n-gram size  
ngram_list =  
list(ngrams(nltk.word_tokenize(sentence), n))  
print(ngram_list)
```

```
[('Generate', 'n-grams'), ('n-grams', 'from'), ('from', 'this'), ('this', 'sentence'), ('sentence', '.')]
```

Named Entity Recognition in Python

```
import spacy
```

```
nlp = spacy.load("en_core_web_sm")  
text = "Apple-Inc.-is-a-technology-company-based-in  
Cupertino,-California."  
doc = nlp(text)
```

```
entities = [(ent.text, ent.label_) for ent in  
doc.ents]  
print(entities)
```

```
[('Apple Inc.', 'ORG'), ('Cupertino', 'GPE'), ('California', 'GPE')]
```